

# CSC2626

## Imitation Learning for Robotics

Florian Shkurti

Week 7: Shared Autonomy and Human-in-the-Loop Learning

# Today's agenda

- Shared autonomy with human in the loop in deep RL
- Hindsight optimization and interactive goal prediction
- Relaxed inverse kinematics for fluid interaction with robot arms


## Acknowledgments

Today's slides are based on student presentations from 2019 by: Andrei Barsan, Bin Yang, and Tingwu Wang

# Shared Autonomy via Deep Reinforcement Learning

Siddharth Reddy, Anca Dragan, Sergey Levine UC Berkeley

Presented by Ioan Andrei Bârsan on February 22, 2019  
[iab@cs.toronto.edu](mailto:iab@cs.toronto.edu)



# Key Question




How can a robot **collaborating** with a human infer the human's goals with as few **assumptions** as possible?

# Motivation


- **Hard:** Actuating a robot with many DoF and/or unfamiliar dynamics.
- **Hard:** Specifying a goal formally (e.g., coordinates).
- **Easy:** Demonstrating the goal indirectly.
  - ...let the machine figure out what I want!



# Motivation: Unknown Dynamics are Hard for Humans



It can get even worse than Lunar Lander...



[www.foddy.net/Athletics.html](http://www.foddy.net/Athletics.html)

or  
Google "qwop"

# Challenges

- **Recall:** Want to demonstrate the goal indirectly with **minimal assumptions**.
  - → We expect the computer to start helping **while it is still learning**.
- **Challenge #1:** How to actually infer user's goal?
- **Challenge #2:** How can we learn this online with low latency?



# **Main Hypothesis**

Shared autonomy can improve human performance without any assumptions about:

(1) dynamics,

(2) the human's policy,

(3) the nature of the goal.

# Formulation: Reward

$$R(s, a, s') = \underbrace{R_{\text{general}}(s, a, s')}_{\text{known}} + \underbrace{R_{\text{feedback}}(s, a, s')}_{\text{unknown, but observed}}$$

↑  
**Agent's reward**  
(what we want to maximize)

↑  
**Handcrafted "common sense"**  
knowledge: do not crash, do  
not tip, etc.

↑  
**Stuff inferred from the human**  
(Main focus of this paper!)

# Formulation

$$\underbrace{R_{\text{feedback}}(s, a, s')}_{\text{unknown, but observed}}$$


- The authors introduce three variants of their method:

Needs  
virtual  
“user”!

1. Known goal space, known user policy.

2. Known goal space, unknown user policy.

3. Unknown goal space, unknown user policy.



**Fewer  
Assumptions**

# The Method

- Based on Q-Learning.
- User input has **two** roles:
  1. A **prior policy** we should fine-tune.
  2. A sensor which can be used to decode the **goal**.
- Short version: Like Q-Learning, but execute closest high-value action to the user's input, instead of highest-value action.

# The Method (Continued)

---

**Algorithm 1** Human-in-the-loop deep Q-learning

---

Standard Q-Learning Initialization

Initialize target action-value function  $Q$  with weights  $\theta^- = \theta$

**for** episode = 1,  $M$  **do**

**for**  $t = 1, T$  **do**

        Sample action  $a_t \sim \pi_\alpha(a_t \mid \tilde{s}_t, a_t^h)$  using equation 3

        Execute action  $a_t$  and observe  $(\tilde{s}_{t+1}, a_{t+1}^h, r_t)$

        Store transition  $(\tilde{s}_t, a_t, r_t, \tilde{s}_{t+1})$  in  $\mathcal{D}$

**if**  $\tilde{s}_{t+1}$  is terminal **then**

**for**  $k = 1$  to  $K$  **do**

▷ training loop

Standard (Double) Q-Learning Training

$\theta \leftarrow \theta - \eta \nabla_\theta \sum_j (y_j - Q(\tilde{s}_j, a_j; \theta))^2$

**end for**

**end if**


    Every  $C$  steps reset  $\hat{Q} = Q$

**end for**

**end for**

---

Interesting part!


$$\pi_\alpha(a \mid \tilde{s}, a^h) = \delta \left( a = \underset{\{a: Q'(\tilde{s}, a) \geq (1-\alpha)Q'(\tilde{s}, a^*)\}}{\arg \max} f(a, a^h) \right),$$

# The Method (Continued)

$$\pi_{\alpha}(a \mid \tilde{s}, a^h) = \delta \left( a = \underset{\{a: Q'(\tilde{s}, a) \geq (1-\alpha)Q'(\tilde{s}, a^*)\}}{\arg \max} f(a, a^h) \right)$$

Maximize similarity to user action

...ensuring action is "close enough" to optimal one.

---

## Algorithm 1 Human-in-the-loop deep Q-learning

---

### Standard Q-Learning Initialization

Initialize target action-value function  $Q'$  with weights  $\theta' = \theta$

**for** episode = 1,  $M$  **do**

**for**  $t = 1, T$  **do**

    Sample action  $a_t \sim \pi_{\alpha}(a_t \mid \tilde{s}_t, a_t^h)$  using equation 3

    Execute action  $a_t$  and observe  $(\tilde{s}_{t+1}, a_{t+1}^h, r_t)$

    Store transition  $(\tilde{s}_t, a_t, r_t, \tilde{s}_{t+1})$  in  $\mathcal{D}$

**if**  $\tilde{s}_{t+1}$  is terminal **then**

**for**  $k = 1$  to  $K$  **do** ▷ training loop

        Sample minibatch  $(\tilde{s}_j, a_j, r_j, \tilde{s}_{j+1})$  from  $\mathcal{D}$

### Standard Training

**end for**

**end if**

    Every  $C$  steps reset  $\hat{Q} = Q$

**end for**

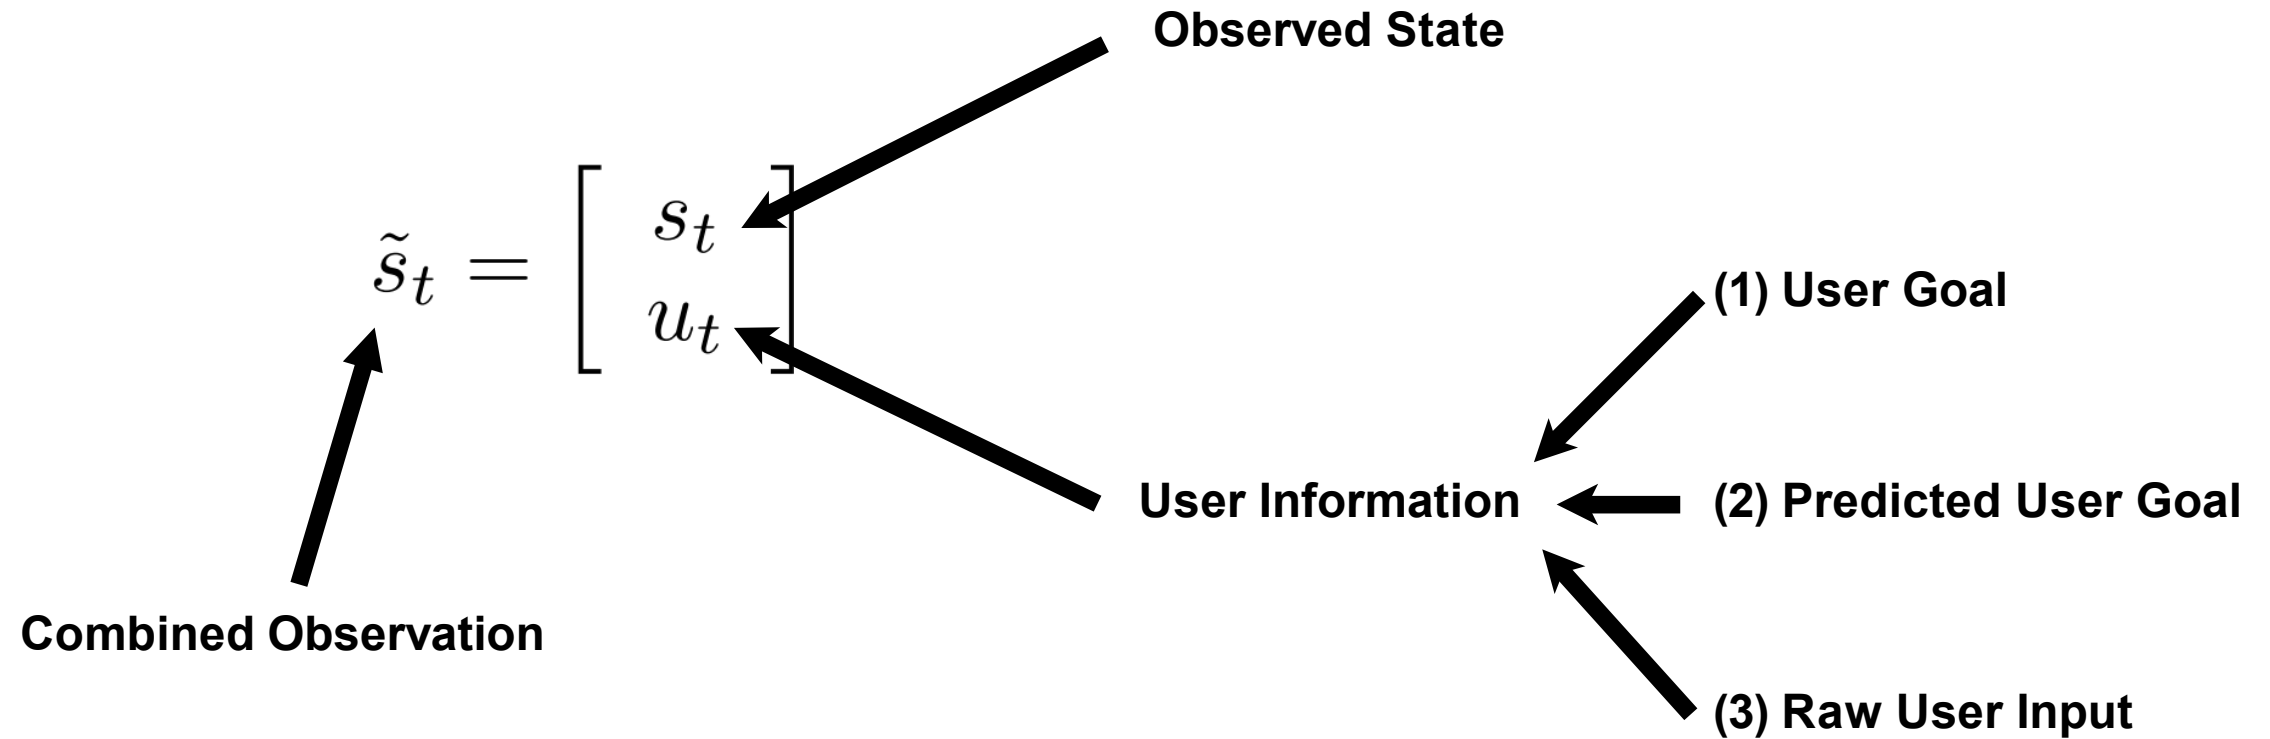
**end for**

---

# But where is $R_{\text{feedback}}$ ?

- The choice of  $R_{\text{feedback}}$  determines what kind of **input** we give to the Q- Learning agent in addition to state!
  1. Known goal space & user policy → exact goal.
  2. Known goal space & unknown policy → predicted goal (pretrained LSTM).
  3. Unknown goal space & policy → the user's input (**main focus**)

# Input to RL Agent






# Experiments


- **Virtual** experiments with Lunar Lander in OpenAI gym.
- **Physical** experiments with an actual drone.

# Real-World Experiments




- **Goal:** Land drone on pad **facing a certain way**.
- **Pilot:** Human, knows target orientation.
- **Copilot:** Our Agent, knows where pad is, but not target orientation.

# Real-World Results



**Important observation: Only  $n = 4$  humans in drone study. 😞**

# Experimental Results: Assumptions



- Higher alpha means we take any action.  $\alpha = 1.0$  means we ignore the pilot.
- Experimented in virtual environment.

# Recap: Strengths

- Good results even when making no assumptions about user/goal.
- Writing is very clear!
- Possible applications in many fields, including e.g., **prosthetics, wheelchairs.**
- Source code released on GitHub!

# Recap: Weaknesses

- User studies could have had more participants.
- Could have shown results on more Gym environments.
- Solution does not generalize to sophisticated long-term goals.


# Conclusion

- Can do shared autonomy with minimal assumptions!
- Idea: Q-Learning & pick high-value action most similar to user's action.
- Works well in virtual environments (real humans).
- Seems to work well in real environments, too.

# Thanks for your attention!

Q&A, if time permits it.

Project website: <https://sites.google.com/view/deep-assist>





**Video of computer-assisted human piloting the lander.**



# Shared Autonomy via Hindsight Optimization

# Teleoperation




**Noisy, insufficient degrees of freedom, tedious**

# Shared Autonomy




**User Input**

+



**Autonomous Assistance**

=



**Achieve Goal**





# Shared Autonomy




**User Input**

+



=



**Achieve Goal**


**Autonomous Assistance**

# Shared Autonomy


## Predict goal Assist for single goal

[Dragan and Srinivasa 13]  
[Kofman et al. 05]  
[Kragic et al. 05]  
[Yu et al. 05]  
[McMullen et al. 14]

...



**Autonomous Assistance**



**Achieve Goal**

# Shared Autonomy

## Predict goal Assist for single goal


[Dragan and Srinivasa 13]  
[Kofman et al. 05]  
[Kragic et al. 05]  
[Yu et al. 05]  
[McMullen et al. 14]  
...

## Predict goal distribution Assist for distribution

[Hauser 13]  
**This work!**




**Autonomous Assistance**



**Achieve Goal**

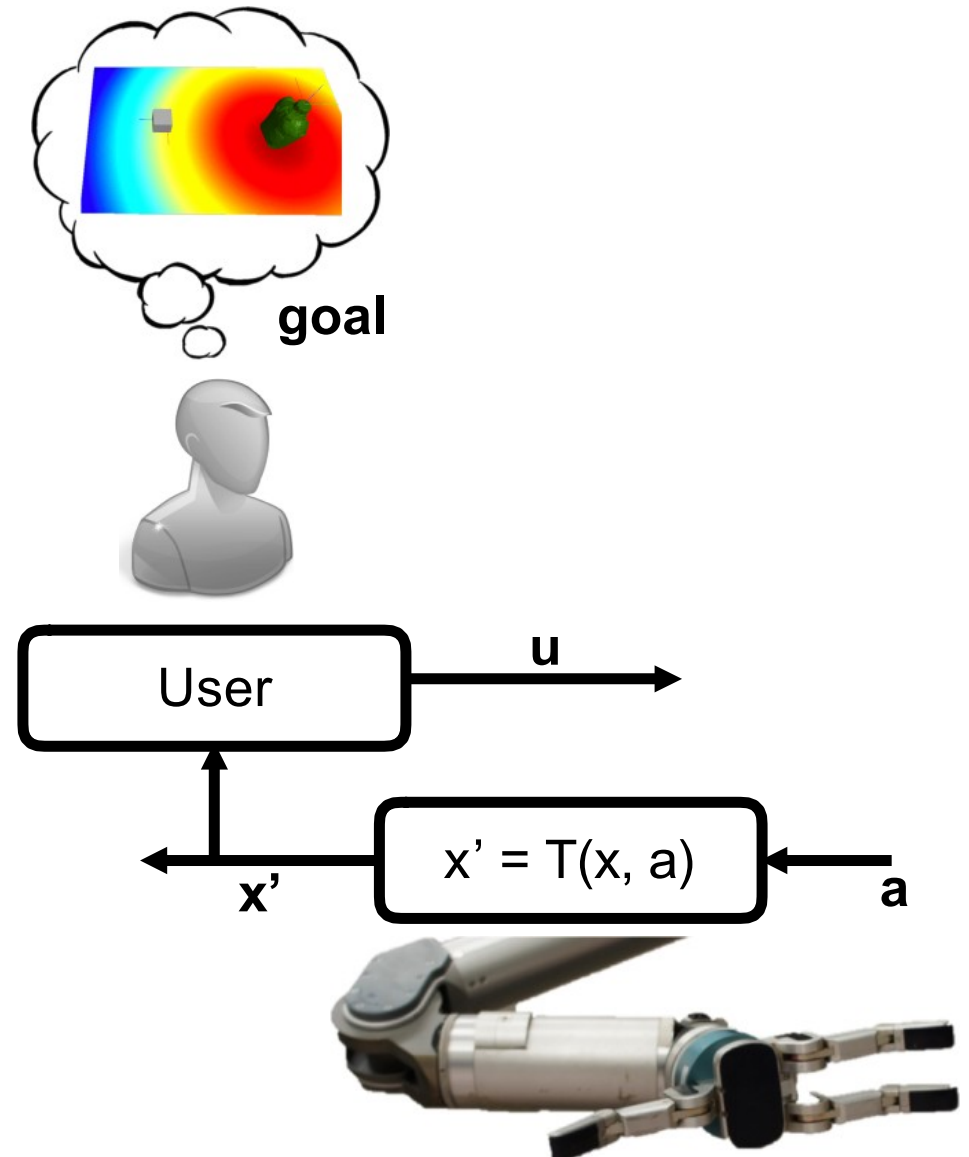
# Method

- System dynamics:  $\mathbf{x}' = T(\mathbf{x}, \mathbf{a})$



# Method

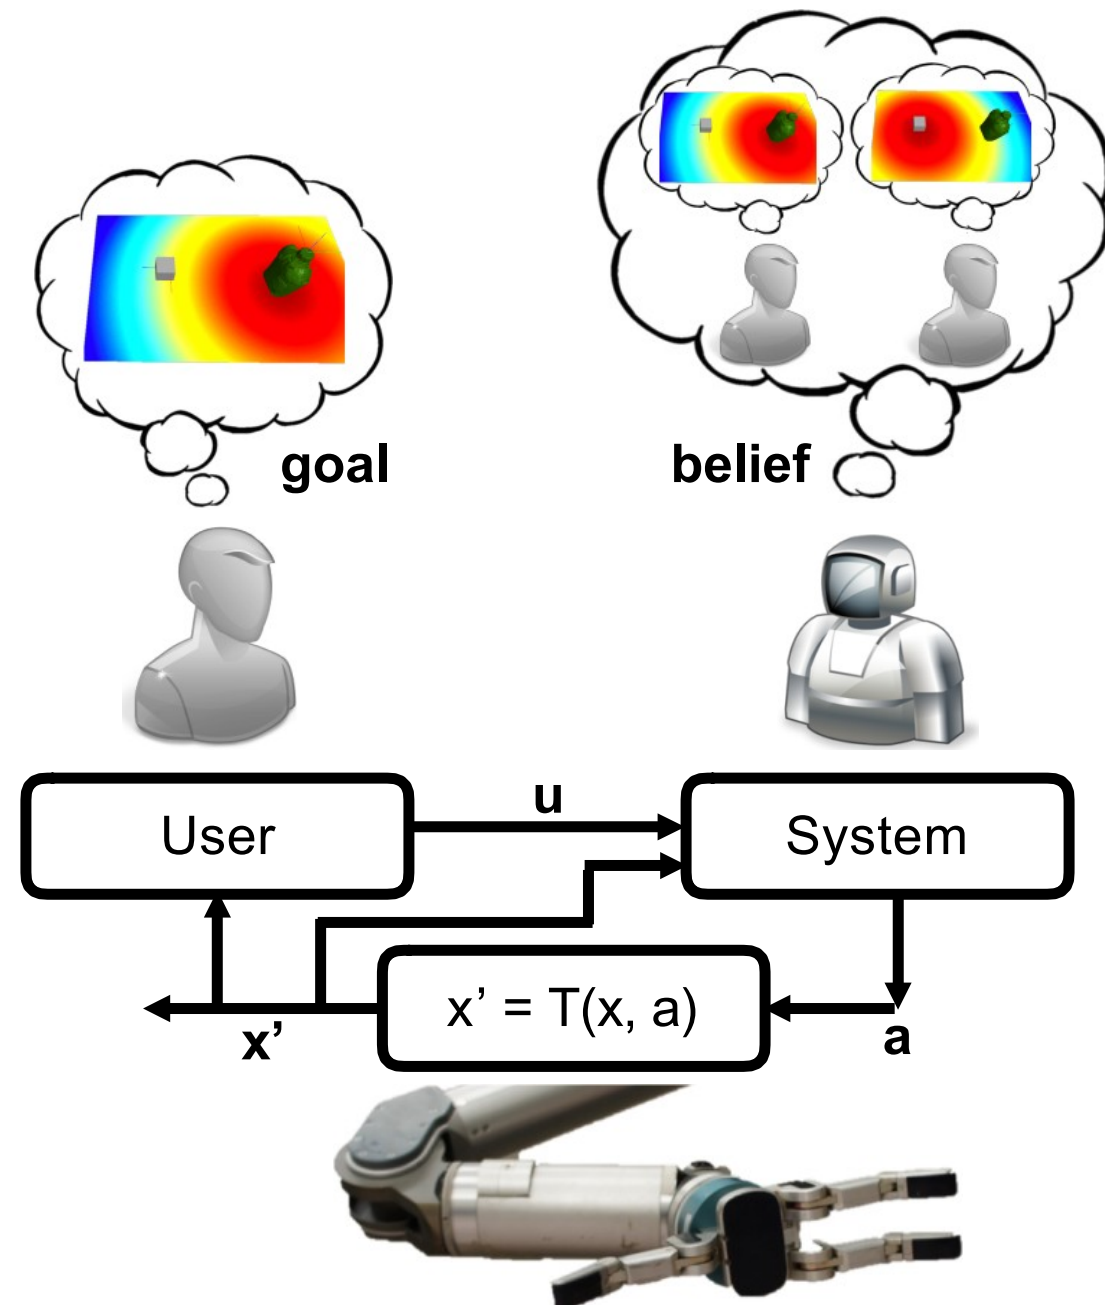
- System dynamics:  $x' = T(x, a)$
- User (MDP) as  $(X, U, T, C_g^{\text{usr}})$ 
  - User policy:  $\pi_g^{\text{usr}}(x) = p(\underline{u} | x, \underline{g})$
  - MaxEnt IOC:  $C_g^{\text{usr}} : X \times U \rightarrow \mathcal{R}$ .





# Method

- System dynamics:  $x' = T(x, a)$
  - User (MDP) as  $(X, U, T, C_g^{\text{usr}})$ 
    - User policy:  $\pi_g^{\text{usr}}(x) = p(u|x, g)$
    - MaxEnt IOC:  $C_g^{\text{usr}} : X \times U \rightarrow \mathcal{R}$ .
  - System (POMDP) as  $(S, A, T, C^{\text{rob}}, U, \Omega)$ 
    - Uncertainty over user's goal
    - System state:  $S = X \times G$
    - Observation: user inputs  $U$
    - Observation model  $\Omega$
- $$p(g|\xi^{0 \rightarrow t}) = \frac{p(\xi^{0 \rightarrow t}|g)p(g)}{\sum_{g'} p(\xi^{0 \rightarrow t}|g')p(g')}$$
- Cost function  $C^{\text{rob}} : S \times A \times U \rightarrow \mathcal{R}$ .



# Hindsight Optimization

- MDP solution:

$$V^{\pi^r}(s) = \mathbb{E} \left[ \sum_t C^r(s_t, u_t, a_t) \mid s_0 = s \right]$$

$$V^*(s) = \min_{\pi^r} V^{\pi^r}(s)$$

# Hindsight Optimization

- MDP solution:

$$V^{\pi^r}(s) = \mathbb{E} \left[ \sum_t C^r(s_t, u_t, a_t) \mid s_0 = s \right]$$

$$V^*(s) = \min_{\pi^r} V^{\pi^r}(s)$$

- POMDP solution:

$$V^{\pi^r}(b) = \mathbb{E} \left[ \sum_t C^r(s_t, u_t, a_t) \mid b_0 = b \right]$$

$$V^*(b) = \min_{\pi^r} V^{\pi^r}(b)$$

# Hindsight Optimization

- MDP solution:

$$V^{\pi^r}(s) = \mathbb{E} \left[ \sum_t C^r(s_t, u_t, a_t) \mid s_0 = s \right]$$

$$V^*(s) = \min_{\pi^r} V^{\pi^r}(s)$$

- POMDP solution:

$$V^{\pi^r}(b) = \mathbb{E} \left[ \sum_t C^r(s_t, u_t, a_t) \mid b_0 = b \right]$$

$$V^*(b) = \min_{\pi^r} V^{\pi^r}(b)$$

- HOP approximation:

$$\begin{aligned} V^{\text{HS}}(b) &= \mathbb{E}_b \left[ \min_{\pi^r} V^{\pi^r}(s) \right] \\ &= \mathbb{E}_g[V_g(x)] \end{aligned}$$

# Hindsight Optimization

- MDP solution:

$$V^{\pi^r}(s) = \mathbb{E} \left[ \sum_t C^r(s_t, u_t, a_t) \mid s_0 = s \right]$$

$$V^*(s) = \min_{\pi^r} V^{\pi^r}(s)$$


- POMDP solution:

$$V^{\pi^r}(b) = \mathbb{E} \left[ \sum_t C^r(s_t, u_t, a_t) \mid b_0 = b \right]$$

$$V^*(b) = \min_{\pi^r} V^{\pi^r}(b)$$


- HOP approximation:

$$\begin{aligned} V^{\text{HS}}(b) &= \mathbb{E}_b \left[ \min_{\pi^r} V^{\pi^r}(s) \right] \\ &= \mathbb{E}_g[V_g(x)] \end{aligned}$$



Deterministic  
problem for  
each future

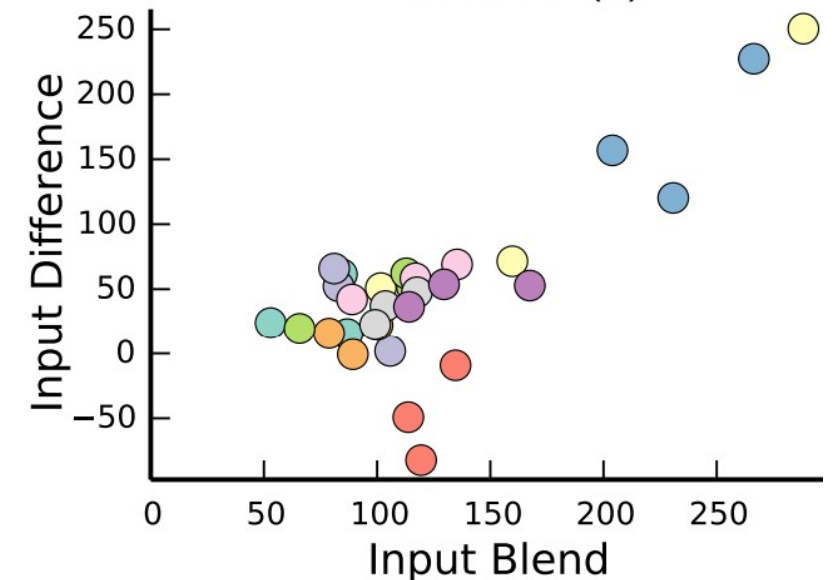

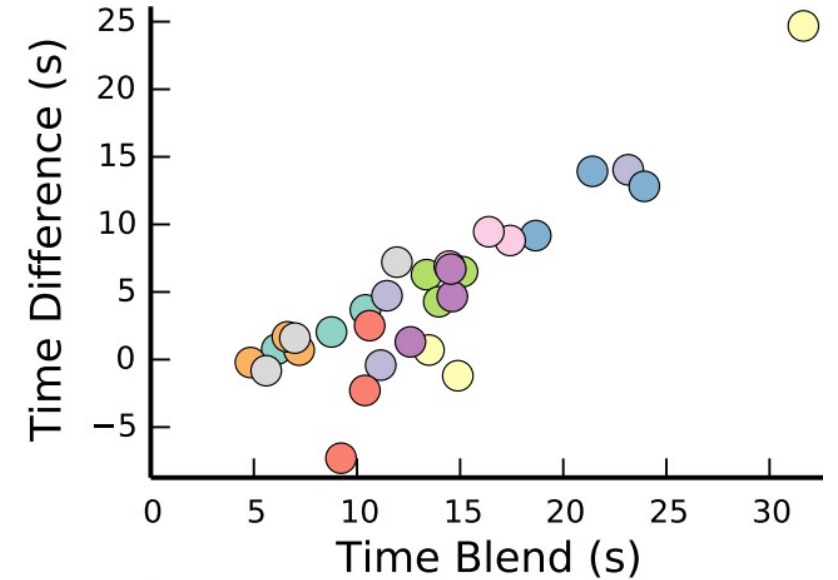

# Results (video)



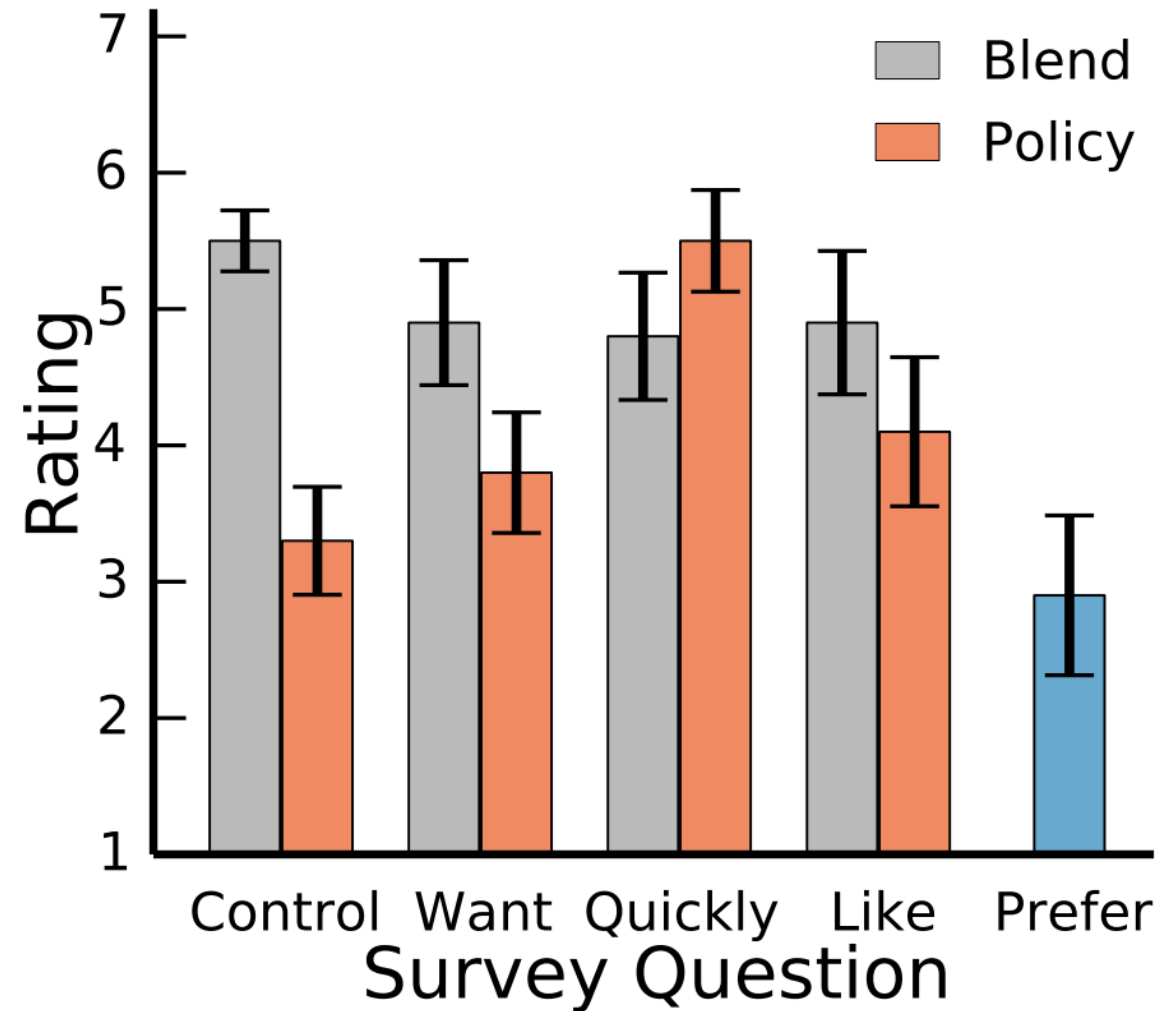
# Results

Compare with method that predicts one goal, the proposed method has:

- Faster execution time
- Fewer user inputs



# User Study





# Limitations

- Requires prior knowledge about the world:
  - a dynamics model that predicts the consequences of taking a given action in a given state of the environment;
  - the set of possible goals for the user;
  - the user's control policy given their goal.
- Suitable in constrained domains where where this knowledge can be directly hard-coded or learned.
- Unsuitable for unstructured environments with ill-defined goals and unpredictable user behavior.

# References

- Javdani, S., Srinivasa, S. S., & Bagnell, J. A. (2015). Shared autonomy via hindsight optimization. *Robotics science and systems: online proceedings, 2015*.
- RSS2015 talk: “Shared autonomy via hindsight optimization”
- Javdani, S., Admoni, H., Pellegrinelli, S., Srinivasa, S. S., & Bagnell, J. A. (2018). Shared autonomy via hindsight optimization for teleoperation and teaming. *The International Journal of Robotics Research*, 37(7), 717-742.
- ICAPS 2015 talk: "Hindsight Optimization for Probabilistic Planning with Factored Actions"



UNIVERSITY OF  
**TORONTO**



VECTOR  
INSTITUTE

INSTITUT  
VECTEUR

# RelaxedIK: Real-time Synthesis of Accurate and Feasible Robot Arm Motion

---




Tingwu Wang

University Of Toronto, CSC-2621, Paper Reading Seminar

# Recap: Forward Kinematics (FK)

---

1. Forward Kinematics
  - a. A common robotic skeleton is a tree of rigid bones
  - b. The relative Euler angles of all the bones determine the end-effectors
    - i. End-effectors? A tool that's connected to the end of a robot arm



# Recap: Inverse Kinematics (IK)

---

## 1. Inverse Kinematics

- a. The indirect control of forward kinematics makes it hard to use in application
  - i. Achieve certain poses?
  - ii. Achieve certain velocities (reward)?
- b. We formulate the inverse kinematics function as:  $\Theta = IK(\mathbf{p})$ , which can be easily written in an analytic form for a simple tree
  - i. Pose contains velocity?
  - ii. Hard to find feasible state space?
- c. In reality, IK is often treated as an optimization problem

$$\chi_p(\Theta) = || \mathbf{p}_g - FK(\Theta) ||_2$$



# Imitation Learning

---


1. Imitation learning has been studied by different communities
  - a. Motion synthesis in character animation?
  - b. Inverse optimal control?
  - c. Imitation learning?

# Imitation Learning

---

1. Imitation learning has been studied by different communities
2. Within the focus of this course, people worked on imitation learning using
  - a. Forward dynamics
    - i. Shooting method (optimize the actions)



Planning via Model Predictive Control



# Imitation Learning

---

1. Imitation learning has been studied by different communities
2. Within the focus of this course, people worked on imitation learning using
  - a. Forward dynamics
  - b. Inverse dynamics
    - i. Collocation method (optimize states)





# Imitation Learning

---

1. Imitation learning has been studied by different communities
2. Within the focus of this course, people worked on imitation learning using
  - a. Forward dynamics
    - i. Shooting method
  - b. Inverse dynamics
    - i. Collocation method
  - c. Model-free method
    - i. GAIL
  - d. Motion synthesis with IK
    - i. Today's paper
    - ii. Old school but with new techniques

# Imitation Learning

---

1. Imitation learning using IK
  - a. Basic idea: Using IK to bridge between target pose and agent's angles
  - b. Input: M (consecutive) expert (goal) poses
  - c. Output: M (consecutive) frames of agent's euler joints
  - d. Constraints:
    - i. IK constraints (goal constraints)
    - ii. Between-frames constraints
    - iii. Etc.

# Imitation Learning


---

1. Imitation learning using IK
  - a. Basic idea: minimize the difference of target pose and agent pose
2. Direct point-to-point approach
  - a. TRAC-IK (previous state-of-the-art)
  - b. Pose2pose / frame2frame imitation learning
    - i. Ignore most of the constraints between frames
  - c. Problems
    - i. Self-collision
      1. Time constraints

# Imitation Learning

---

## 1. Self-collision



# Imitation Learning

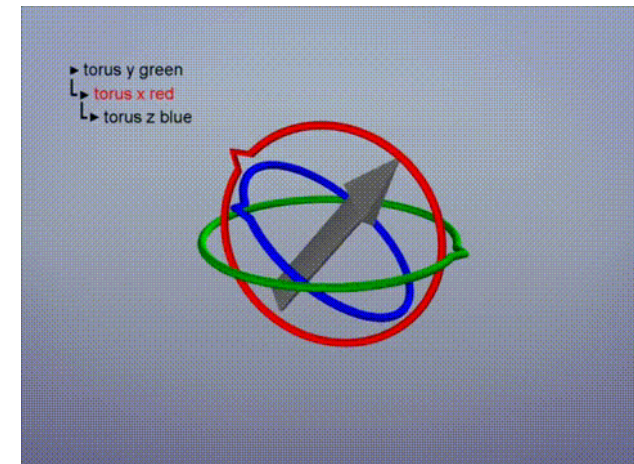
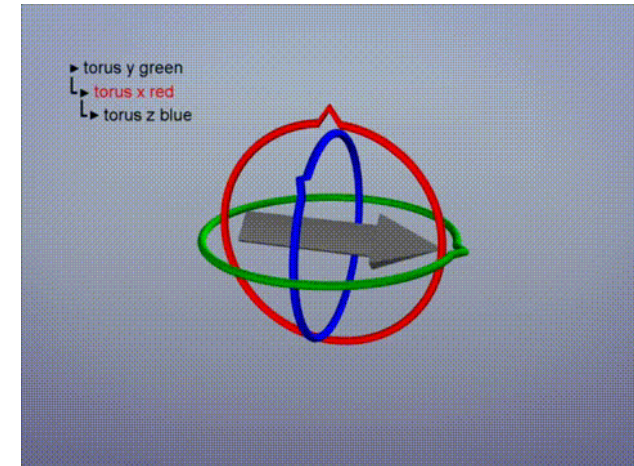

---

1. Imitation learning using IK
  - a. Basic idea: minimize the difference of target pose and agent pose
2. Direct point-to-point approach
  - a. TRAC-IK (previous state-of-the-art)
  - b. Pose2pose / frame2frame imitation learning
    - i. Ignore most of the constraints between frames
  - c. Problems
    - i. Self-collision
    - ii. Singularities



# Imitation Learning

1. Singularities
  - a. E.g. Losing a DoF
  - b. Infinite control signals



# Imitation Learning

---


1. Imitation learning using IK
  - a. Basic idea: minimize the difference of target pose and agent pose
2. Direct point-to-point approach
  - a. TRAC-IK (previous state-of-the-art)
  - b. Pose2pose / frame2frame imitation learning
    - i. Ignore most of the constraints between frames
  - c. Problems
    - i. Self-collision
    - ii. Singularities
    - iii. Discontinuity



# Imitation Learning

---

## 1. Discontinuity





# Imitation Learning

---

1. Imitation learning using IK
  - a. Basic idea: minimize the difference of target pose and agent pose
2. Direct point-to-point approach
3. Real-time motion planning approach
  - a. Output the (conservative) solutions real-time
    - i. Always meet the control & collision constraints
    - ii. Soft goal constraints
  - b. Problems
    - i. Goal mistracking

# Imitation Learning

---

## 1. Goal mistracking



# Imitation Learning


---

1. Imitation learning using IK
  - a. Basic idea: minimize the difference of target pose and agent pose
2. Direct point-to-point approach
3. Real-time motion planning approach
  - a. Output the (conservative) solutions real-time
    - i. Always meet the control & collision constraints
    - ii. Soft goal constraints
  - b. Problems
    - i. Goal mistracking
    - ii. Unpredictable behaviors

# Imitation Learning

---

1. Unpredictable behaviors





# Relaxed IK

---

1. Basic Idea: Using soft (Relaxed) IK loss that considers self-collision and singularity for faster optimization

$$\mathbf{f}(\Theta) = \sum_{i=1}^k w_i f_i(\Theta, \Omega_i)$$


1. Loss functions
  - a. End-effector position & orientation matching
  - b. Minimize joint velocity, acceleration, jerk
  - c. Self-collision loss (fast)
  - d. Singularity loss

# Relaxed IK

---

1. Self-collision loss
  - a. Common approach: very slow
  - b. Relaxed IK:
    - i. Approximate how imminent the robot is to a collision state
    - ii. Using simulated data to train a network to predict the distances between links

$$col(\Theta) = \sum_{i,j} b * exp(\frac{-dis(l_i, l_j)^2}{2c^2})$$



# Relaxed IK

---

1. Singularity loss
  - a. Kinematic singularities are well studied in robotics
  - b. Relaxed IK:
    - i. Find a metric that can approximate distance to a singularity
    - ii. Jacobian's condition number is used as a proxy distance to singularity
      1. Why?

$$\dot{\mathbf{x}} = \mathbf{J}(\Theta)\dot{\Theta}$$


- i. Penalize condition values less than mean – b \* std
  1. Estimate mean, std from simulated data

# Relaxed IK

---

## 1. Pros

- a. Much faster and smoother performance
  - i. combining neural network and traditional robotics
- b. Data driven, less human-engineering
  - i. Novel singularity metric
- c. Easy to deploy
  - i. Sim2Real





# Relaxed IK

---

1. Pros
2. Cons
  - a. No safety / convergence guarantee
  - b. Weak experiments section
    - i. Under-tuned baseline
    - ii. Limited ablation study
  - c. Slower than point2point methods
  - d. Hyper-parameter sensitive

